Message from the Program Chairs

ICCV 2013

Welcome to Sydney, Australia and the 14th IEEE International Conference on Computer Vision (ICCV). In addition to the main four day program of oral and poster presentations, ICCV 2014 has 25 co-located workshops, 10 tutorials, and demos and exhibits.

We received 1629 submissions to the conference (a record for ICCV). Of these, 41 (2.5%) were accepted as orals, and 413 (25.4%) as posters, for an overall acceptance rate of 27.9% See http://www.iccv2013.org/decisions.php for a more detailed breakdown of acceptance rates as a function of research area and review scores.

To handle the review process, we invited 50 leading vision researchers to act as Areas Chairs (ACs) and recruited 1152 expert reviewers from the computer vision community, with a maximum of 11 papers per reviewer (most had far fewer). Our candidate reviewer list consisted of the CVPR 2013 reviewers, with the addition of area chairs from the last several computer vision conferences and editorial board members from leading vision journals. We were extremely impressed with the diligence of the reviewers across the board. Out of nearly 5000 reviews (each paper had three reviews), only a couple dozen were not returned—an extraordinarily small number. For these papers, we called on a set of senior "emergency reviewers," not in the regular reviewer pool, who graciously agreed in advance to handle one or two papers at the last minute with a day or two turn-around time.

We used the CMT conference management service sponsored by Microsoft Research to manage the submission and selection of papers from beginning to end. The AC meeting was facilitated by a set of new scripts to give detailed and individualized per-AC and per-panel paper/review summaries distributed via dropbox.

After the submission deadline, we asked ACs to bid on preferred papers to review, and combined the resulting bids with the automated Toronto Paper Matching System (TPMS) to assign papers to ACs. TPMS, develped by Charlin et al. [UAI 2011], suggests matches between papers and reviewers (ACs, in our case) based on bag-of-words descriptors extracted from the PDF files of submitted manuscripts and representative publications by each potential reviewer. The ACs in turn used the results of a TPMS matching of papers to reviewers to help them determine the potential reviewers for each of their assigned papers, from which the CMT system automatically selected three non-conflicted reviewers per paper.

Conflicts were handled differently this year. In the past, if you had a collaborator at (e.g.,) MIT, you had to declare a conflict with the entire institution, thereby ruling out all reviewers at MIT. Furthermore, the need to self-declare conflicts led to omissions, as many authors failed to list all their collaborator's institutions. Instead, this year, we instructed authors to only list their own institutions (not collaborators' institutions), and we auto-detected collaborators using an online publication database (DBLP). For each author A, we marked a reviewer R as a conflict if A and R co-authored a paper in the last 3 years, and one of them was first author. In addition, reviewers could notify us in case they noticed (or suspected) conflicts. However, we received only a handful of these, indicating that the auto-detection strategy worked well.

Reviewers were given four weeks to complete their reviews, at which time the ACs stepped back in to vet the reviews for quality (initiating discussions with reviewers, where necessary) before they were released to the authors. After the author rebuttals were collected, a second discussion phase ensued.

New this year, we asked reviewers to rate papers both before AND after the author's rebuttal. Furthermore, in the latter (post-rebuttal) rating, we removed the borderline option, and instead required

reviewers to commit to recommending the paper as reject, (weak) poster, or (weak) oral. While this required more work on the part or reviewers, the vast majority completed both ratings, and the ACs found them extremely useful in decision making.

The AC meeting was held at Oxford, to make decisions on all but the clear rejects. Clear reject decisions were made by the individual AC handling each paper, and verified by a second AC. Roughly two thirds of submissions made it to the meeting. On the first day of the AC meeting, every paper was discussed by an AC triple, with the goal of making accept/reject decisions. Each AC was part of two different triples, one that met in the morning, and the other in the afternoon, to provide more fine-grained handling of conflicts. Only papers not rejected at the end of day one made it to day two. Day two focused on poster/oral decisions, and all oral candidates were discussed by a panel of about a dozen ACs. Each AC was part of two different panels, one that met in the morning, and the other in the afternoon (again, for more fine-grained conflict handling). The Program Chairs served as the panel chairs. By the end of the meeting, the ACs were asked to produce detailed consolidation reports to justify all their decisions.

The Program Chairs and General Chairs did not submit any papers to ICCV 2014, allowing them to work without any direct conflicts throughout the review process. Additionally, the ACs were excluded from any decisions associated with papers that they authored or from their affiliated institutions or close collaborators. Triples had no conflicts, and AC panels had papers authored by the ACs in that panel. Soft (institutional or collaborator) conflicts were allowed in panels, with all conflicted ACs leaving the room when each paper was discussed. The double-blind nature of the CVPR review process was thus strictly maintained throughout the review process.

ICCV 2014 is co-sponsored by IEEE and The Computer Vision Foundation (CVF). The proceedings of ICCV 2014 are being published in USB drive form. All papers in the main conference and associated workshops will be indexed by the IEEE, and available through the IEEE Computer Society Digital Library and under IEEE Xplore.

We wish to thank the other members of the Organizing Committee, the Area Chairs, Reviewers, Authors, and the CMT team for the immense amount of hard work and professionalism that has gone into making ICCV 2014. Our thanks also go to the organizers of CVPR 2013 for their helpful advice and support. We are grateful to the sponsors as well. Finally, we wish all the attendees a stimulating, informative, and enjoyable conference.

Kyros Kutulakos, Yi Ma, Steve Seitz, and Phil Torr
Program Co-chairs